

Scaling ML-KEM to (much) higher security

Swechchha Adhikari¹, Vadim Lyubashevsky², and Peter Schwabe^{3,4}

¹ Brigham Young University, Provo, USA
adhikar6@student.byu.edu

² IBM Research, Zurich, Switzerland
VAD@zurich.ibm.com

³ Max Planck Institute for Security and Privacy, Bochum, Germany
peter@cryptojedi.org

⁴ Radboud University, Nijmegen, The Netherlands

1 Introduction

After a multi-year, multi-round public competition, NIST published its first post-quantum standards FIPS-203, FIPS-204, and FIPS-205 in 2024, specifying ML-KEM, ML-DSA, and SLH-DSA respectively. Interestingly, four of the schemes in the PQC Round 3 Submissions Kyber (standardized as ML-KEM), Saber, NTRU LPRime, and FrodoKEM are closely related lattice KEMs, spanning a spectrum of algebraic structure, from FrodoKEM's unstructured LWE, through Kyber's and Saber's cyclotomic modules, to the prime-degree number fields of NTRU Prime. This reflects different design principles, commonly understood as structure vs conservatism. ML-KEM's module structure gives small keys and fast NTT arithmetic, but is also seen as a potential attack surface, so conservative designs discard it entirely. FrodoKEM, for example, is based on a more conservative hardness assumption of plain LWE, which is why bodies such as BSI and ANSSI point to FrodoKEM for long-term confidentiality; it was also recently standardized by ISO for this very reason.

It is useful to separate the two notions of conservatism this framing bundles together as they concern two independent (and today, entirely hypothetical) classes of attack. The first is the choice of assumption. FrodoKEM rests on unstructured LWE to protect against a structure break, an attack that would exploit the algebraic structure of ring/module LWE, but not apply to plain LWE. The second is the security margin a scheme carries above its target level, which protects against advances in the best known attacks today. This could potentially include improved lattice reduction, dual attacks, or sharper cost models that would affect the concrete security of every lattice scheme at its current parameters, including FrodoKEM. Both of these hedges are distinct and cannot be substituted: structured conservatism does nothing against the second class of attacks and no security margin protects against the first. On top of that, it's impossible to predict which of the two threats is more likely.

Faced with this uncertainty, increasing the margin has been a historically standard response. For computationally secure schemes, concrete security estimate is the cost of the best attack known at the time, and such estimates evolve

as cryptanalysis matures. For instance, the recommended size of an RSA modulus climbed from 512 to 1024 to 2048 bits and beyond as factoring algorithms and available computation advanced. Lattice estimates are not static either. The cost models used for estimation and the attacks have evolved over more than a decade, and the cost of the best known attacks continues to drift in both directions as improvements are proposed, scrutinized, and refined.

If the best known attacks continue to improve, how much margin is required remains open, and the useful question is not whether today’s KEMs carry enough but how much they can carry at all. Margin is not free: it is purchased with larger parameters, and hence with bandwidth and runtime, so no scheme can absorb an unbounded sequence of cryptanalytic improvements simply by scaling its parameters. For LWE-family KEMs there is a further constraint, since the parameters that determine security also control the decryption-failure probability, so increasing the noise to gain a larger security margin necessarily affects correctness.

Crucially, what structure determines is the cost of buying additional margin. At matched security levels, an unstructured scheme already requires roughly an order of magnitude more bandwidth than a module scheme. Since this bandwidth regime is considered practical for conservative deployments, a module scheme can spend the same bandwidth budget on larger parameters and obtain a substantially higher security margin while retaining the efficiency advantages of NTT arithmetic. This motivates inverting the usual question. Rather than asking how small a module KEM can be at a fixed security level, we ask how much security margin a module KEM can achieve within the bandwidth regime of a conservative unstructured scheme.

We answer this question with KAIBURR, a module KEM obtained by scaling ML-KEM’s module rank and adjusting the noise distribution to preserve correctness. KAIBURR achieves large security margins within FrodoKEM’s communication cost while running substantially faster.

2 Preliminaries

3 Warmup: Just changing k

4 Revisiting the noise distribution

5 Changing q

6 Benchmarks

Acknowledgements

Large parts of the work were done while the first author was employed at MPI-SP.